

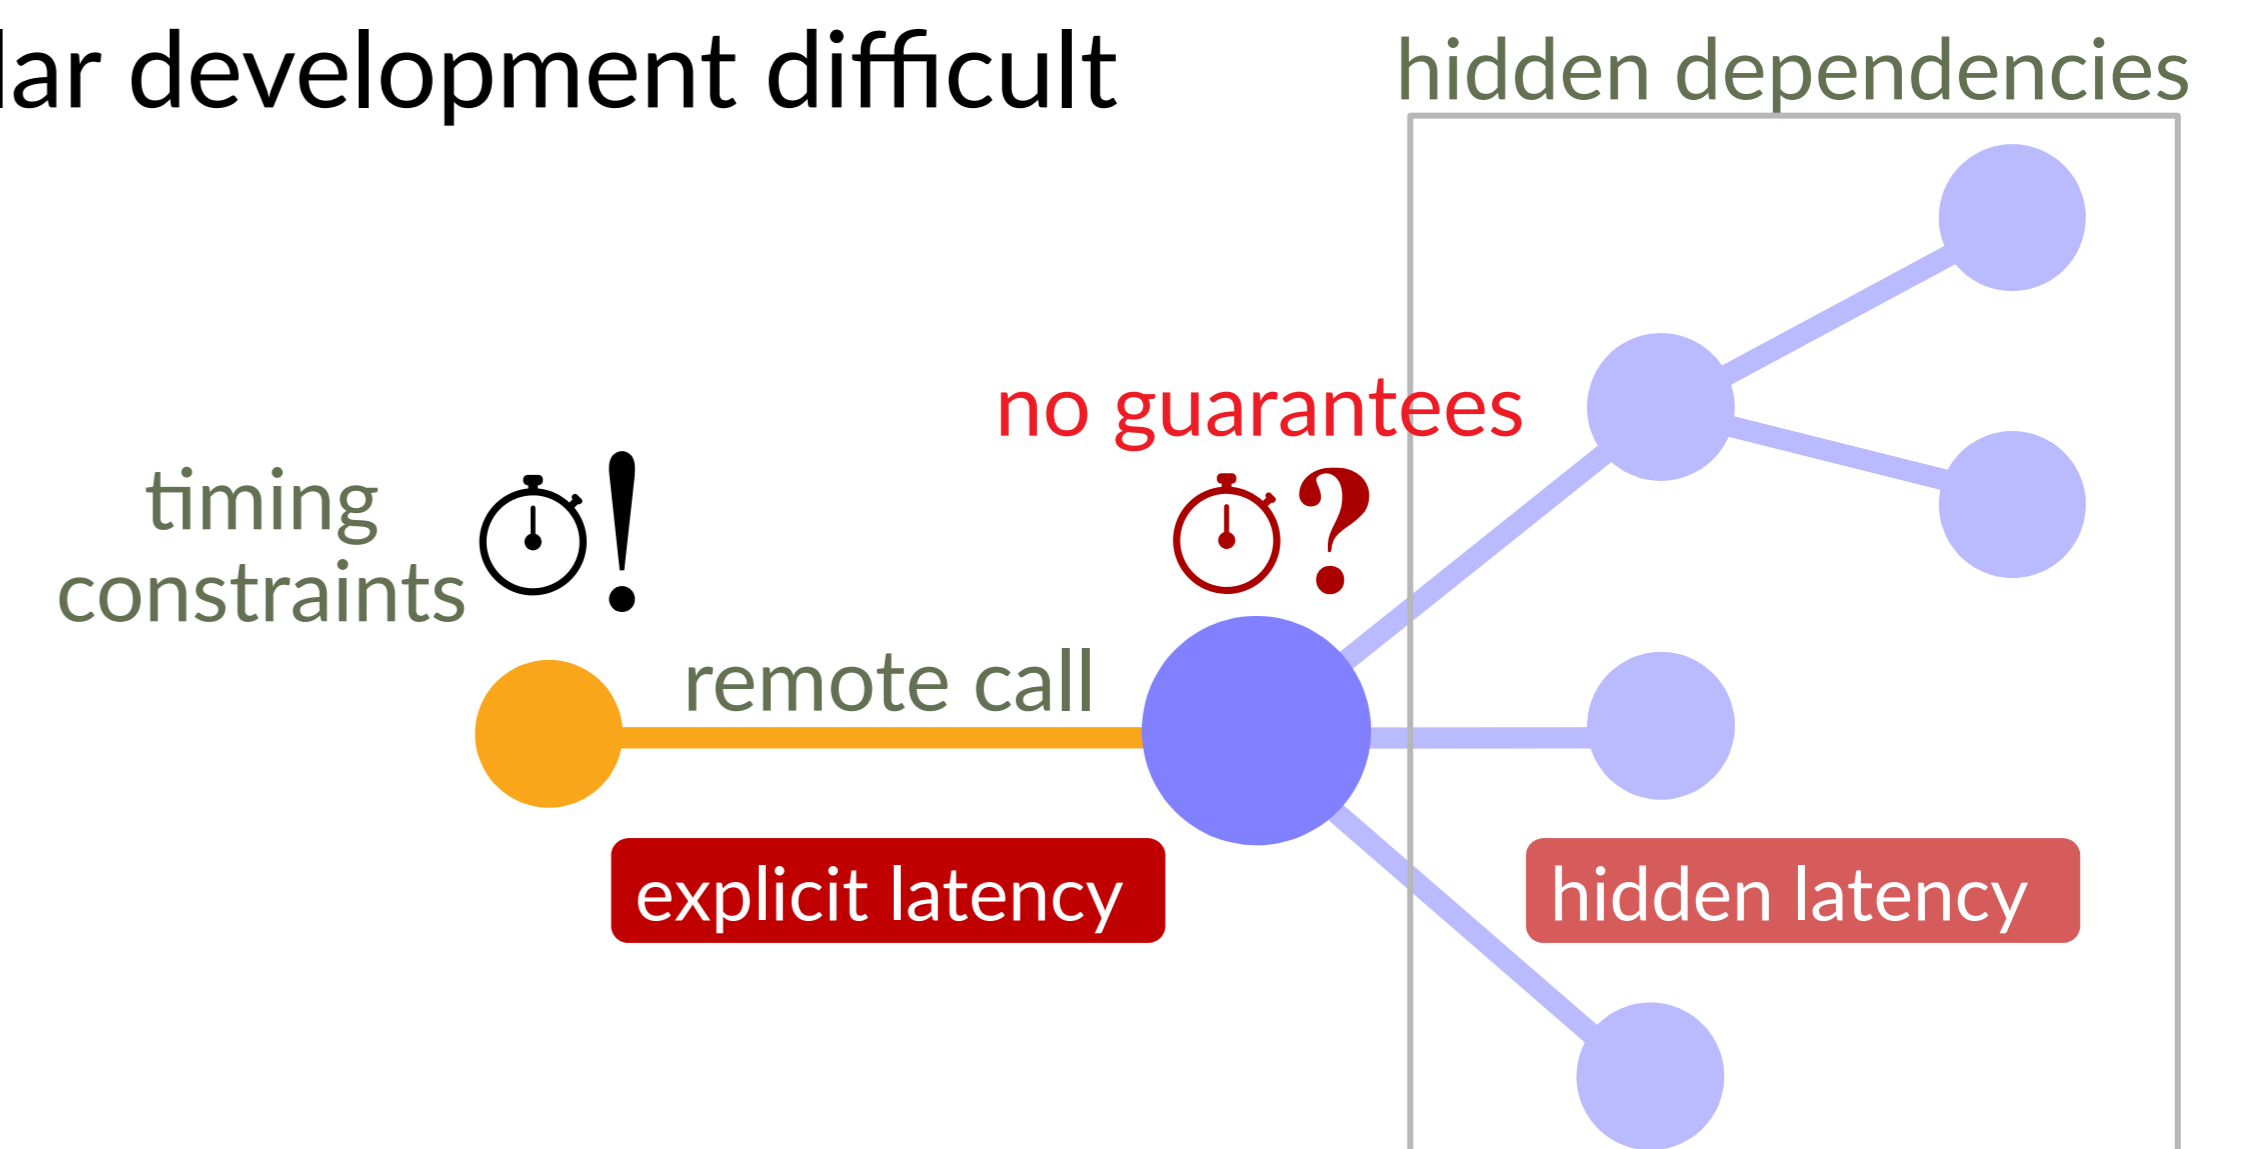
Software across Geo-Distributed Data Centers

- Fixed locations
- Predictable latency between servers:
 - within same data center < 2ms
 - in different data centers > 100ms
- Location matters



Latency is Hard to Predict in Distributed Applications


- Remote computations with hidden dependencies
- No static guarantees towards timing constraints
- Need global view to reason about latency
- Modular development difficult



Goals

Make latency and locations explicit

Static guarantees towards latency



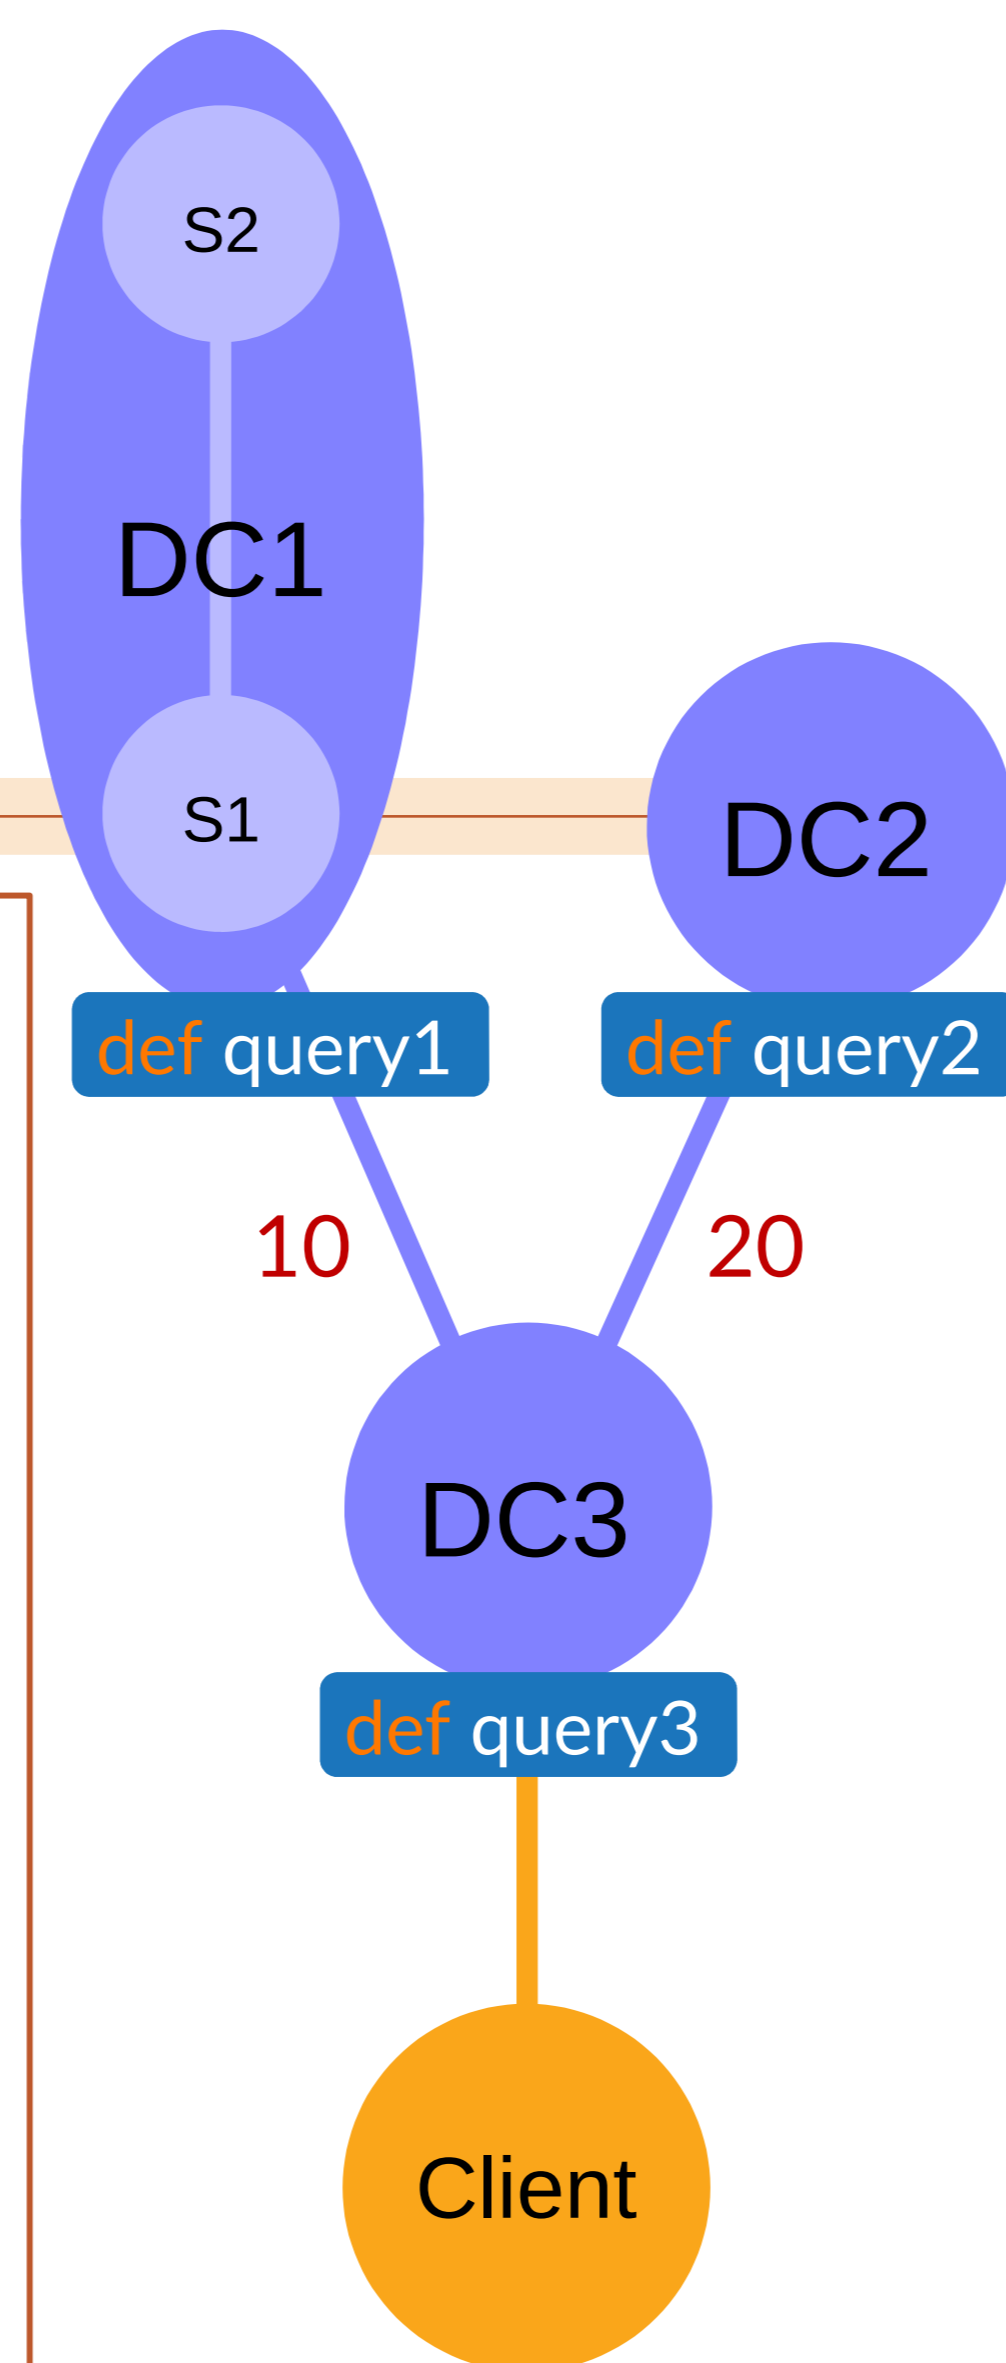
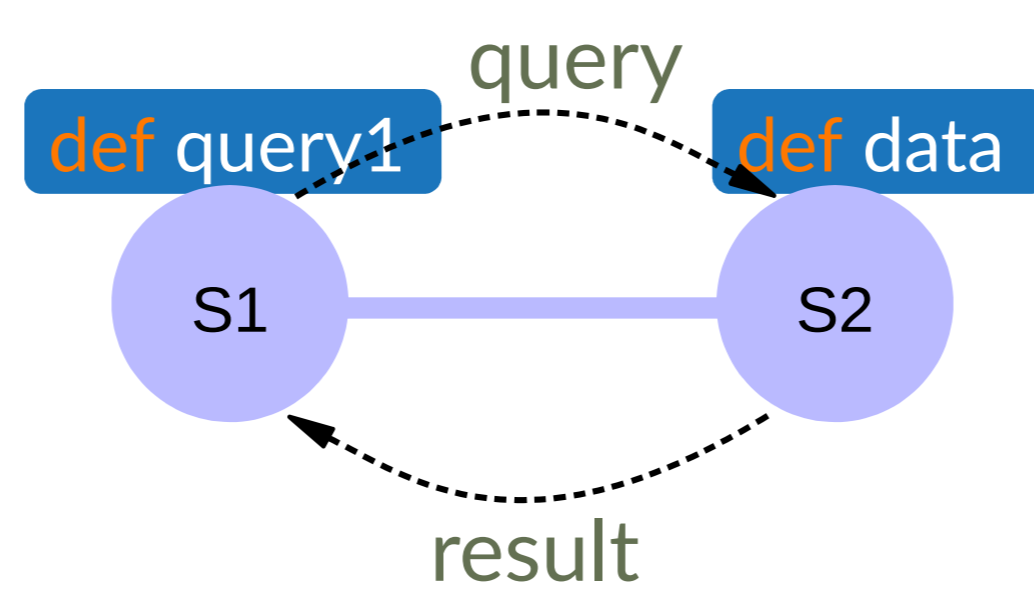
Allows latency-aware software development & refactoring

Extending ScalaLoc's Placement with Latency

- Types encode location and number of remote communication steps.
- Makes locations and latency explicit.
- No hidden dependencies

```
def data(q: Query): Result on S2 withLat 0
  = placed[S2] { computeResult(q) }
  # messages

def query1(q: Query): Result on S1 withLat 2
  = placed[S1] { remoteCall data(q) }
```



Latency Weights & Bounds

- Weights on connections approximate latency predictions.
 - DC1 ↔ DC3 : 10
 - DC2 ↔ DC3 : 20
- Latency is upper bound on all paths through the program.

```
def query3(q: Query): Result on DC3 withLat Max[_40, _22]
  = placed[DC3] {
    if (cond(q)) remoteCall query2(q) latency: 40
    else remoteCall query1(q) latency: 20 + 2
  }
```

Static Guarantees

Type system rejects wrong assumptions:

```
placed[Client] { val r: Result withLat 2
  = remoteCall query3(myQuery); ... }
  latency: 2 + 40
```

Allows explicit over-approximation:

```
placed[Client] { val r: Result withLat 50
  = remoteCall query3(myQuery); ... }
```

Type signatures encode location and latency:

```
query3 : Query => Result on DC3 withLat 40
```

→ Allows modular development.

Provably Correct Bounds

Formalization based on λ -calculus with

- Remote communication primitives
- Placement & latency types
- Sized types & size-decreasing recursion

$$\frac{\Delta; \Gamma; \Lambda; P \vdash t_c : (B_t, [s_t], [l_t]) \quad \Delta; \Gamma; \Lambda; P \vdash t_f : (B_f, [s_f], [l_f])}{\Delta; \Gamma; \Lambda; P \vdash \text{if } t_c \{t_t\} \{t_f\} : (B_t, [s_t], [l_c + \max(l_t, l_f)])} \text{ (T-If)}$$

→ Correctness proof for inferred latency bounds

$$\begin{aligned} & \Delta; \Gamma; \Lambda; P \vdash t : (B, [s], [l_T]) \\ & \bullet (\langle t \rangle_{\mathcal{I}}, [0]) \stackrel{\mathcal{I}}{\rightsquigarrow}_* (\langle v \rangle_{\mathcal{I}}, [l_R]) \implies l_R \leq l_T \end{aligned}$$

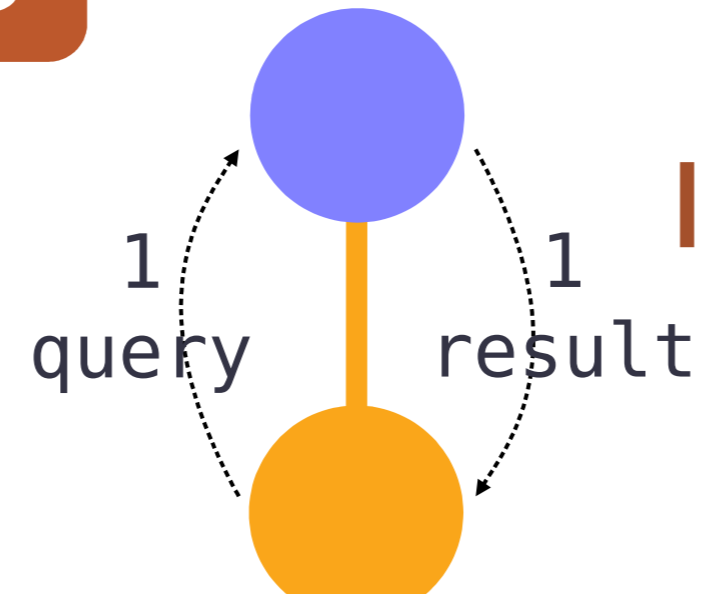
Latency-Saving Refactorings

Type system can guide refactoring.

```
val queries: SizedList[Query, N] = ...
queries.map(q => remoteCall query3(q))
```

N requests

lat. overhead: $N * 2$



Exploit locality to reduce communication.

```
val queries: SizedList[Query, N] = ...
placedCall[DC3] queries.map(q => query3(q))
```

on DC3

→ Can be automated based on type info.

1 request

lat. overhead: 2

